

航空出行管理系统数据库优化报告

小组成员：郭梦琪 张晗翀 刘勉之 刘书畅

一、业务背景描述

随着我国经济的高速发展，国民生活水平的不断提高，人们的出行方式的要求越来越高，利用飞机出行已经成为了不少人的首选，飞机出行有方便、快捷、舒适的特点。对于航空公司来说，业务量逐步增加，需要管理上百万的用户，上千万订单记录，上百的航班等。过多的数据会减缓增删改查的速度，为了保持良好的用户体验，提升管理效率，需要对航空出行管理系统进行优化，将查询时间控制在可控范围内，并需要考虑到多用户并发查询的情况。

二、优化目标

- 降低单人查询时间，包括但不限于查询订单、航班、飞行记录、登机牌信息等
- 提高并发查询速度
- 使用查询缓存，再次查询可提高速度（加分项）

三、性能测试与分析

1. 查询优化

本项目的查询功能较多，我们均进行了优化，重点分析较为典型的查询订单函数queryOrder。

· 查询订单queryOrder

① 函数优化

查询订单函数需要返回航班信息，登机牌信息等多项信息，而这些信息分别储存在boardingpass、flight、route、order多个表里，需要进行join操作。并且用户通常是通过自己的用户id而非order表里，并非主键查询。

优化前的函数：

```
select `order`.orderId,time,boardId,flightId,number,seat,type,price
from `order` natural join boardingpass
where `order`.CustomerId=CustomerId
```

针对该查询的特点我们做了以下优化：

- 增加索引：
boardingpass: flightId CustomerId
order: CustomerId
route: departureTime departureAirport arrivalAirport
- 将flight表的status改为enum类型，因为固定内容范围的查询会比varchar快

3. 若语句使用join加上 `order.orderId=boardingpass.orderId`和 `order.CustomerId=CustomerId`条件，是order这个数据量较小的表作为驱动表，boardingpass这个数据量很大的表作为被驱动表，检索的行数较少时间较少；若语句使用join加上 `order.orderId=boardingpass.orderId`和 `boarding.CustomerId=CustomerId`条件或者使用 `natural join`，是在数据量大的表里进行检索，检索的行数很多，时间较慢
4. 替换select* 按需取数据

优化后的函数：

```
(select `order`.orderId,time,boardId,flightId,number,seat,type,price
from `order` join boardingpass
where `order`.orderId=boardingpass.orderId and `order`.CustomerId=1)
```

② 利用mysqlslap进行性能测试

*注：此测试设置了两组数据：大数据组有500000条boardingpass数据，小数据组有5000条boardingpass数据。部分测试极高并发下情况使用小数据组进行测试。

此数据针对测试函数设计，用户1的登机牌数和订单数非常多，采用用户1进行测试，所以用户1的查询时间比正常情况下要高。


10用户并发 迭代10次（大数据）

优化前：

```
C:\Users\97537>mysqlslap --no-defaults -u root --create-schema=old_large_data --query="call queryOrder(1)" -c 100 -i 5 -p586970
mysqlslap: [Warning] Using a password on the command line interface can be insecure.
C:\Users\97537>mysqlslap --no-defaults -u root --create-schema=old_large_data --query="call queryOrder(1)" -c 10 -i 5 -p586970
mysqlslap: [Warning] Using a password on the command line interface can be insecure.
Benchmark
  Average number of seconds to run all queries: 12.087 seconds
  Minimum number of seconds to run all queries: 9.640 seconds
  Maximum number of seconds to run all queries: 13.735 seconds
  Number of clients running queries: 10
  Average number of queries per client: 1
```

优化后：

```
C:\Users\97537>mysqlslap --no-defaults -u root --create-schema=new_large_data --query="call queryOrder(1)" -c 10 -i 5 -p586970
mysqlslap: [Warning] Using a password on the command line interface can be insecure.
Benchmark
  Average number of seconds to run all queries: 0.012 seconds
  Minimum number of seconds to run all queries: 0.000 seconds
  Maximum number of seconds to run all queries: 0.032 seconds
  Number of clients running queries: 10
  Average number of queries per client: 1
```



③ 利用explain分析性能

优化前：

| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
|----|-------------|----------|------------|--------|------------------------|---------|---------|-----------|--------|----------|-------------|
| 1 | SIMPLE | boarding | (Null) | ref | FK_belongOrder,FK_take | FK_take | 4 | const | 248956 | 100 | Using index |
| 1 | SIMPLE | order | (Null) | eq_ref | PRIMARY,FK_buy | PRIMARY | 4 | old_large | 1 | 50 | Using where |

优化后：

| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
|----|-------------|----------|------------|------|----------------------|----------------|---------|-----------|-------|----------|--------|
| 1 | SIMPLE | order | (Null) | ref | PRIMARY,FK_buy,order | FK_buy | 4 | const | 25305 | 100 | (Null) |
| 1 | SIMPLE | boarding | (Null) | ref | FK_belongOrder | FK_belongOrder | 4 | new_large | 14 | 100 | (Null) |

对比possible_keys，可以看到优化后的函数使用了新建的索引。

对比rows可以发现优化后的函数检索行数大大减少。

· queryFlight 查询航班

①函数优化

1. 由于查询条件是根据日期和起飞降落地的，所以给departureAirport, arrivalAirport建立了索引
2. 由于如果将某个信息设为允许null会增加存储查询成本，所以在还不知道actualDeparture actualArrival时间时，填为scheduleDeparture scheduleArrival
3. 由于exists的查询效率比in高，将关键字in 改为了exists

优化前函数：

```

create
    definer = root@localhost procedure queryFlight(IN date varchar(20), IN
departureAirport varchar(30),
                                                    IN arrivalAirprot
varchar(30))
select *
from flight
where flight.routeId in (select routeId from route where
DATE(route.departureTime)=date and route.departureAirport=departureAirport
and route.arrivalAirport=arrivalAirport)
;

```

优化后函数：

```

create
    definer = root@localhost procedure queryFlight(IN date varchar(20), IN
departureAirport varchar(30),
                                                    IN arrivalAirprot
varchar(30))
begin
    select *
    from flight
    where exists(select routeId from route where
DATE(route.departureTime)=date and route.departureAirport=departureAirport
and route.arrivalAirport=arrivalAirport and route.routeId=flight.routeId);
end;

```

② 利用mysqlslap进行性能测试

优化前：

```

C:\Users\97537>mysqlslap --no-defaults -u root --create-schema=old_large_data --query="call queryFlight('2018
-09-10','cjkfghjklbnm', fghjklvbnmrtjkff)" -c 500 -i 10 -p586970
mysqlslap: [Warning] Using a password on the command line interface can be insecure.
Benchmark
Average number of seconds to run all queries: 0.953 seconds
Minimum number of seconds to run all queries: 0.141 seconds
Maximum number of seconds to run all queries: 2.156 seconds
Number of clients running queries: 500
Average number of queries per client: 1

```

优化后：

```

C:\Users\97537>mysqlslap --no-defaults -u root --create-schema=new_large_data --query="call queryFlight('
-09-10','cjkfghjklbnm', fghjklvbnmrtjkff)" -c 500 -i 10 -p586970
mysqlslap: [Warning] Using a password on the command line interface can be insecure.
Benchmark
Average number of seconds to run all queries: 0.709 seconds
Minimum number of seconds to run all queries: 0.187 seconds
Maximum number of seconds to run all queries: 1.765 seconds
Number of clients running queries: 500
Average number of queries per client: 1

```

③ 利用explain分析性能

优化前：

| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
|----|-------------|--------|------------|------|----------------|---------|---------|----------|------|----------|-------------|
| 1 | SIMPLE | route | (Null) | ALL | PRIMARY | (Null) | (Null) | (Null) | 50 | | 2 Using whe |
| 1 | SIMPLE | flight | (Null) | ref | FK_belongRoute | FK_belo | 4 | airplane | 10 | 100 | (Null) |

优化后：

| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
|----|--------------------|--------|------------|--------|--------------------------------|---------|---------|-----------|------|----------|-------------|
| 1 | PRIMARY | flight | (Null) | ALL | (Null) | (Null) | (Null) | (Null) | 500 | 100 | Using where |
| 2 | DEPENDENT SUBQUERY | route | (Null) | eq_ref | PRIMARY,departureAirport_index | PRIMARY | 4 | airplane1 | 5 | 5 | Using where |

对比possible_keys, 可以看到优化后的函数使用了新建的索引, 自建的索引发挥了作用。

对比type,可以看到优化后提升为eq_ref

*注: 其他查询函数queryUnfinishedTake、queryFinishTake和queryFlight大同小异, 就不再赘述。

2. 查询缓存

① 什么是查询缓存:

MySQL查询缓存保存查询返回的完整结果。当查询命中该缓存, MySQL会直接返回结果, 跳过了解析、优化和执行截断。

查询缓存系统会跟踪查询中涉及的每个表, 如果这些表发生变化, 那么和这个表相关的所有的缓存数据都将失效。这种机制效率看起来比较低, 因为数据表变化时很有可能对应的查询结果没有变更, 但是这种简单实现代价很小, 而这点对于一个非常繁忙的系统来说非常重要。

随着现在的通用服务器越来越强大, 查询缓存被发现是一个影响服务器扩展性的因素。它可能成为整个服务器的资源竞争单点, 在多核服务器上还可能导致服务器僵死。所以大部分时候应该默认关闭查询缓存, 如果查询缓存作用很大的话, 可以配置个几十兆的小缓存空间。

② 使用方法

1. 在mysql安装目录下的my.ini(windows)中进行修改, 在其中加入“query_cache_type=1”后重启mysql可启用缓存
2. SHOW VARIABLES LIKE '%query_cache%';可检查当前查询缓存相关配置参数
3. 参数解释:

query_cache_type: 查询缓存类型,是否打开缓存

可选项

- a、0(OFF): 关闭 Query Cache 功能, 任何情况下都不会使用 Query Cache;
- b、1(ON): 开启 Query Cache 功能, 但是当SELECT语句中使用SQL_NO_CACHE提示后, 将不使用Query Cache;
- c、2(DEMAND): 开启Query Cache 功能, 但是只有当SELECT语句中使用了SQL_CACHE 提示后, 才使用Query Cache。

备注1:

如果query_cache_type为on而又不想利用查询缓存中的数据, 可以用下面的SQL:

```
SELECT SQL_NO_CACHE * FROM my_table WHERE condition;
```

如果值为2, 要使用缓存的话, 需要使用SQL_CACHE开关参数:

```
SELECT SQL_CACHE * FROM my_table WHERE condition;
```

query_cache_size: 缓存使用的总内存空间大小,单位是字节,这个值必须是1024的整数倍,否则MySQL 会自动调整降低最小量以达到1024的倍数; (感觉这个应该跟文件系统的block大小有关)

query_cache_min_res_unit: 分配内存块时的最小单位大小, 设置查询缓存Query Cache每次分配内存的最小空间大小, 即每个查询的缓存最小占用的内存空间大小;

query_cache_limit: 允许缓存的单条查询结果集的最大容量，默认是1MB，超过此参数设置的查询结果集将不会被缓存；

query_cache_wlock_invalidate: 如果某个数据表被锁住,是否仍然从缓存中返回数据,默认是OFF,表示仍然可以返回

控制当有写锁定发生在表上的时刻是否先失效该表相关的Query Cache，如果设置为 1(TRUE)，则在写锁定的同时将失效该表相关的所有Query Cache，如果设置为0(FALSE)则在锁定时刻仍然允许读取该表相关的Query Cache。

③ 测试配置：

| Variable_name | Value |
|--------------------|-----------|
| have_query_cache | YES |
| query_cache_limit | 268435456 |
| query_cache_min_re | 4096 |
| query_cache_size | 268435456 |
| query_cache_type | ON |
| query_cache_wlock_ | OFF |

④ 测试结果

开启缓存后多次测试时间对比

```
C:\Users\liumi>mysqlslap --no-defaults -u root --create-schema=gmqairplane --query="select * from (select * from boardin
gpass where boardingpass.CustomerId=1) as B,(select * from flight natural join route) as FR where (select F.actualArriva
l from flight as F where B.flightId=F.flightId)<FR.actualDeparture;" -c 5 i 10
Benchmark
Average number of seconds to run all queries: 22.313 seconds
Minimum number of seconds to run all queries: 22.313 seconds
Maximum number of seconds to run all queries: 22.313 seconds
Number of clients running queries: 5
Average number of queries per client: 1

C:\Users\liumi>mysqlslap --no-defaults -u root --create-schema=gmqairplane --query="select * from (select * from boardin
gpass where boardingpass.CustomerId=2) as B,(select * from flight natural join route) as FR where (select F.actualArriva
l from flight as F where B.flightId=F.flightId)<FR.actualDeparture;" -c 5 i 10
Benchmark
Average number of seconds to run all queries: 0.015 seconds
Minimum number of seconds to run all queries: 0.015 seconds
Maximum number of seconds to run all queries: 0.015 seconds
Number of clients running queries: 5
Average number of queries per client: 1
```

对比结果可以看到，打开缓存后第一次操作时间需要22.313秒，而第二次测试时速度明显提升，只需要0.015秒，缓存大大提升了查询的速度。

值得注意的是，call function无法触发查询缓存机制，其不产生cache_inserts，原因未知。在理想的实际使用中，应是包装好的软件直接使用sql语句而非调用函数，故仍然可用。

3.数据库存储引擎

① MYSQL5.7中的引擎

| Engine | Support | Comment | Transactions | XA | Savepoints |
|--------------------|---------|--|--------------|-----|------------|
| InnoDB | DEFAULT | Supports transactions, row-level locking | YES | YES | YES |
| MRG_MYISAM | YES | Collection of identical MyISAM tables | NO | NO | NO |
| MEMORY | YES | Hash based, stored in memory | NO | NO | NO |
| BLACKHOLE | YES | /dev/null storage engine | NO | NO | NO |
| MyISAM | YES | MyISAM storage engine | NO | NO | NO |
| CSV | YES | CSV storage engine | NO | NO | NO |
| ARCHIVE | YES | Archive storage engine | NO | NO | NO |
| PERFORMANCE_SCHEMA | YES | Performance Schema | NO | NO | NO |

InnoDB存储引擎介绍

MySQL版本>=5.5 默认的存储引擎，MySQL推荐使用的存储引擎。支持事务，行级锁定，外键约束。事务安全型存储引擎。更加注重数据的完整性和安全性。

MyISAM存储引擎介绍

MySQL<= 5.5 MySQL默认的存储引擎。

ISAM: Indexed Sequential Access Method (索引顺序) 的缩写，是一种文件系统。

擅长与处理，高速读与写。

myisam支持全文索引，innodb在5.6之后支持

myisam可以用myisamPack压缩数据，节省磁盘空间，缺点是压缩后需要重新修复索引且变为只读，修改需要重新解压。非常适用于route和过往记录等极少修改的数据的储存

关于Innodb 和myisam的取舍：

Innodb：数据完整性，并发性处理，擅长更新，删除。

myisam：高速查询及插入。擅长插入和查询，**不支持事务**。

② 测试

```
C:\Users\liumi>mysqlslap --no-defaults -u root -a -c 1 i 100 --engine=myisam,innodb
Benchmark
Running for engine myisam
Average number of seconds to run all queries: 0.219 seconds
Minimum number of seconds to run all queries: 0.219 seconds
Maximum number of seconds to run all queries: 0.219 seconds
Number of clients running queries: 1
Average number of queries per client: 0

Benchmark
Running for engine innodb
Average number of seconds to run all queries: 0.390 seconds
Minimum number of seconds to run all queries: 0.390 seconds
Maximum number of seconds to run all queries: 0.390 seconds
Number of clients running queries: 1
Average number of queries per client: 0
```



```
C:\Users\liumi>mysqlslap --no-defaults -u root --create-schema=gmqairplane --query="call queryOrder(1)" -c 10 i 5 --engine=myisam,innodb
Benchmark
Running for engine myisam
Average number of seconds to run all queries: 33.828 seconds
Minimum number of seconds to run all queries: 33.828 seconds
Maximum number of seconds to run all queries: 33.828 seconds
Number of clients running queries: 10
Average number of queries per client: 1

Benchmark
Running for engine innodb
Average number of seconds to run all queries: 35.766 seconds
Minimum number of seconds to run all queries: 35.766 seconds
Maximum number of seconds to run all queries: 35.766 seconds
Number of clients running queries: 10
Average number of queries per client: 1
```

由上图测试结果可以发现，在低并发性的时候两者的效率相近，且myisam始终略优。

但是在高并发的情况下innodb对myisam体现出显著优势：

```
C:\Users\liumi>mysqlslap --no-defaults -u root -a -c 100 i 100 --engine=myisam,innodb
Benchmark
Running for engine myisam
Average number of seconds to run all queries: 51.453 seconds
Minimum number of seconds to run all queries: 51.453 seconds
Maximum number of seconds to run all queries: 51.453 seconds
Number of clients running queries: 100
Average number of queries per client: 0

Benchmark
Running for engine innodb
Average number of seconds to run all queries: 5.500 seconds
Minimum number of seconds to run all queries: 5.500 seconds
Maximum number of seconds to run all queries: 5.500 seconds
Number of clients running queries: 100
Average number of queries per client: 0
```

③对于我们的航空出行管理系统，如何选择

1. 由于航空出行管理系统需要经常进行数据的修改和删除，InnoDB较为适合
2. 在进行变更时，如果失败需要回滚必须用到事务，InnoDB较为适合
3. 每个用户账户数据的完整性和同步性非常重要，需要外键支持，否则会导致混乱，InnoDB较为适合。

四、总结

查询优化：

1. 建立索引
2. 将部分范围固定的数据类型从varchar改为enum
3. 用exists代替in
4. 减少join操作，并先进行select再join，减少join操作需要操作的函数
5. 把所有字段都改为not null

查询缓存：

多次相同或者相近的查询则会由于cache_hits，速度大大加快

存储引擎：

低并发性的时候两者的效率相近，且myisam始终略优

但是在高并发的情况下innodb对myisam体现出显著优势

且由于需要用到事务，支持高并发，经常修改，此系统更适合InnoDB

测试：

最后进行汇总对比。

在缓存属性为如下图的情况下：

| | |
|-----|---|
| | 1000 0000 |
| HEX | 1000 0000 |
| DEC | 268,435,456 |
| OCT | 2 000 000 000 |
| BIN | 0001 0000 0000 0000 0000 0000 0000 0000 |

优化前无缓存+两种引擎对比：

```
C:\Users\liumi>mysqlslap --no-defaults -u root --create-schema=airplane_old_large --query="select `order`.orderId,time,boardId,flightId,number,seat,type,price from `order` join boardingpass where `order`.orderId=boardingpass.orderId and `order`.CustomerId=1;" -c 50 i 5 --engine=mysam,innodb
Benchmark
Running for engine mysam
Average number of seconds to run all queries: 31.454 seconds
Minimum number of seconds to run all queries: 31.454 seconds
Maximum number of seconds to run all queries: 31.454 seconds
Number of clients running queries: 50
Average number of queries per client: 1

Benchmark
Running for engine innodb
Average number of seconds to run all queries: 20.484 seconds
Minimum number of seconds to run all queries: 20.484 seconds
Maximum number of seconds to run all queries: 20.484 seconds
Number of clients running queries: 50
Average number of queries per client: 1
```

优化前有缓存+两种引擎对比：

```
C:\Users\liumi>mysqlslap --no-defaults -u root --create-schema=airplane_old_large --query="select `order`.orderId,time,boardId,flightId,number,seat,type,price from `order` join boardingpass where `order`.orderId=boardingpass.orderId and `order`.CustomerId=1;" -c 50 i 5 --engine=mysam,innodb
Benchmark
Running for engine mysam
Average number of seconds to run all queries: 5.156 seconds
Minimum number of seconds to run all queries: 5.156 seconds
Maximum number of seconds to run all queries: 5.156 seconds
Number of clients running queries: 50
Average number of queries per client: 1

Benchmark
Running for engine innodb
Average number of seconds to run all queries: 4.203 seconds
Minimum number of seconds to run all queries: 4.203 seconds
Maximum number of seconds to run all queries: 4.203 seconds
Number of clients running queries: 50
Average number of queries per client: 1
```

优化后有缓存+两种引擎对比：

```
C:\Users\liumi>mysqlslap --no-defaults -u root --create-schema=airplane_new_large --query="select t1.orderId,time,boardId,flightId,number,seat,type,price from (select orderId,time from `order` where (`order`.CustomerId=1)) t1, (select boardId,orderId,flightId,number,seat,type,price from boardingpass where boardingpass.CustomerId=1) t2 where t1.orderId=t2.orderId;" -c 50 i 5 --engine=mysam,innodb
Benchmark
Running for engine mysam
Average number of seconds to run all queries: 3.813 seconds
Minimum number of seconds to run all queries: 3.813 seconds
Maximum number of seconds to run all queries: 3.813 seconds
Number of clients running queries: 50
Average number of queries per client: 1

Benchmark
Running for engine innodb
Average number of seconds to run all queries: 3.875 seconds
Minimum number of seconds to run all queries: 3.875 seconds
Maximum number of seconds to run all queries: 3.875 seconds
Number of clients running queries: 50
Average number of queries per client: 1
```

可以看到针对我们的航空出行管理系统，进行优化和开启缓存，并使用innodb引擎，可以即达提升性能。

综合对比图

